

# **GPU101 Project**

## **Breadth-First Search Algorithm**

Lorenzo Di Napoli

[lorenzo.dinapoli@mail.polimi.it](mailto:lorenzo.dinapoli@mail.polimi.it)

# Project Goal

## Breadth-First Search Algorithm

- Implementation and optimisation of the ***Breadth-First Search (BFS)*** Algorithm in *CUDA* programming language for the *GPU*.

# Starting Materials


- a .mtx file (size 317x317) containing the Tree Structure.
- an already working C++ script wich:
  - reads and represents the .mtx file in CSR format (*Compressed Sparse Row*).
  - executes via *CPU* the ***BFS\_sequential()*** Algorithm.

# How to Read the Tree

- In order to graphically visualise the .mtx file, I wrote a python script.
- The script works this way: from an .mtx input file it displays a graphic of the matrix.

```
1  import matplotlib.pyplot as plt
2  import numpy as np
3  import os
4
5  # Check if the file exists
6  file_path = 'cavity02.mtx'
7  if not os.path.exists(file_path):
8      print(f"File {file_path} not found.")
9      exit(1)
10
11  try:
12      # Read the matrix manually
13      with open(file_path, 'r') as f:
14          lines = f.readlines()
15
16      # Skip comments and read the header
17      lines = [line for line in lines if not line.startswith('%')]
18      header = lines[0].strip().split()
19      num_rows, num_cols, num_entries = map(int, header)
20
21      # Initialize the matrix
22      matrix = np.zeros((num_rows, num_cols))
23
24      # Read the entries
25      for line in lines[1:]:
26          row, col, value = map(float, line.strip().split())
27          matrix[int(row)-1, int(col)-1] = value
28
29      # Print matrix details for debugging
30      print(f"Matrix type: {type(matrix)}")
31      print(f"Matrix shape: {matrix.shape}")
32
33      # Plot the matrix
34      plt.spy(matrix, markersize=1)
35      plt.title('Visualization of ' + file_path)
36      plt.show()
37  except Exception as e:
38      print(f"An error occurred: {e}")
```



# Visualization of the Tree



# Explanation of the Algorithm

## BFS vs DFS

- **Breadth-First Search (BFS)** is a traversing algorithm which works by exploring the totality of the frontier before going deeper (Breadth analysis).
- It is different from the **Depth-First Search (DFS)**, which works by exploring first by depth and then by frontier (Depth analysis).



# Procedure

## main( ) adptation

- In order to call the GPU-kernell, it is necessary to adapt the **main( )** function:
  - have to be initialised some pointers for the variables to give at the kernell, paying attention to the double cases using *host\_variable* / *gpu\_variable*.
  - have to be allocated the *GPU* device memory via `cudaMalloc( )` and copied the data in the direction *HostToDevice* via `cudaMemcpy( )`, after the kernell had executed the *BFS*, the memory is freed via `cudaFree( )`.
  - all the *CUDA* calls are checked via the macro `CHECK( call )`, the kernel launch is checked via the macro `CHECK_KERNELLCALL( )`.



# Procedure

## main( ) adptation

```
// gpu variables allocation
int *gpu_row_ptr;
int *gpu_col_ind;
int *gpu_dist;
int *gpu_source;

// gpu memory allocation
CHECK(cudaMalloc(&gpu_source, sizeof(int)));
CHECK(cudaMalloc(&gpu_row_ptr, host_row_ptr.size() * sizeof(int)));
CHECK(cudaMalloc(&gpu_col_ind, host_col_ind.size() * sizeof(int)));
CHECK(cudaMalloc(&gpu_dist, num_vals * sizeof(int)));

// Copy data from host to device
CHECK(cudaMemcpy(gpu_source, &host_source, sizeof(int), cudaMemcpyHostToDevice));
CHECK(cudaMemcpy(gpu_row_ptr, host_row_ptr.data(), host_row_ptr.size() * sizeof(int), cudaMemcpyHostToDevice));
CHECK(cudaMemcpy(gpu_col_ind, host_col_ind.data(), host_col_ind.size() * sizeof(int), cudaMemcpyHostToDevice));
CHECK(cudaMemcpy(gpu_dist, host_dist.data(), num_vals * sizeof(int), cudaMemcpyHostToDevice));

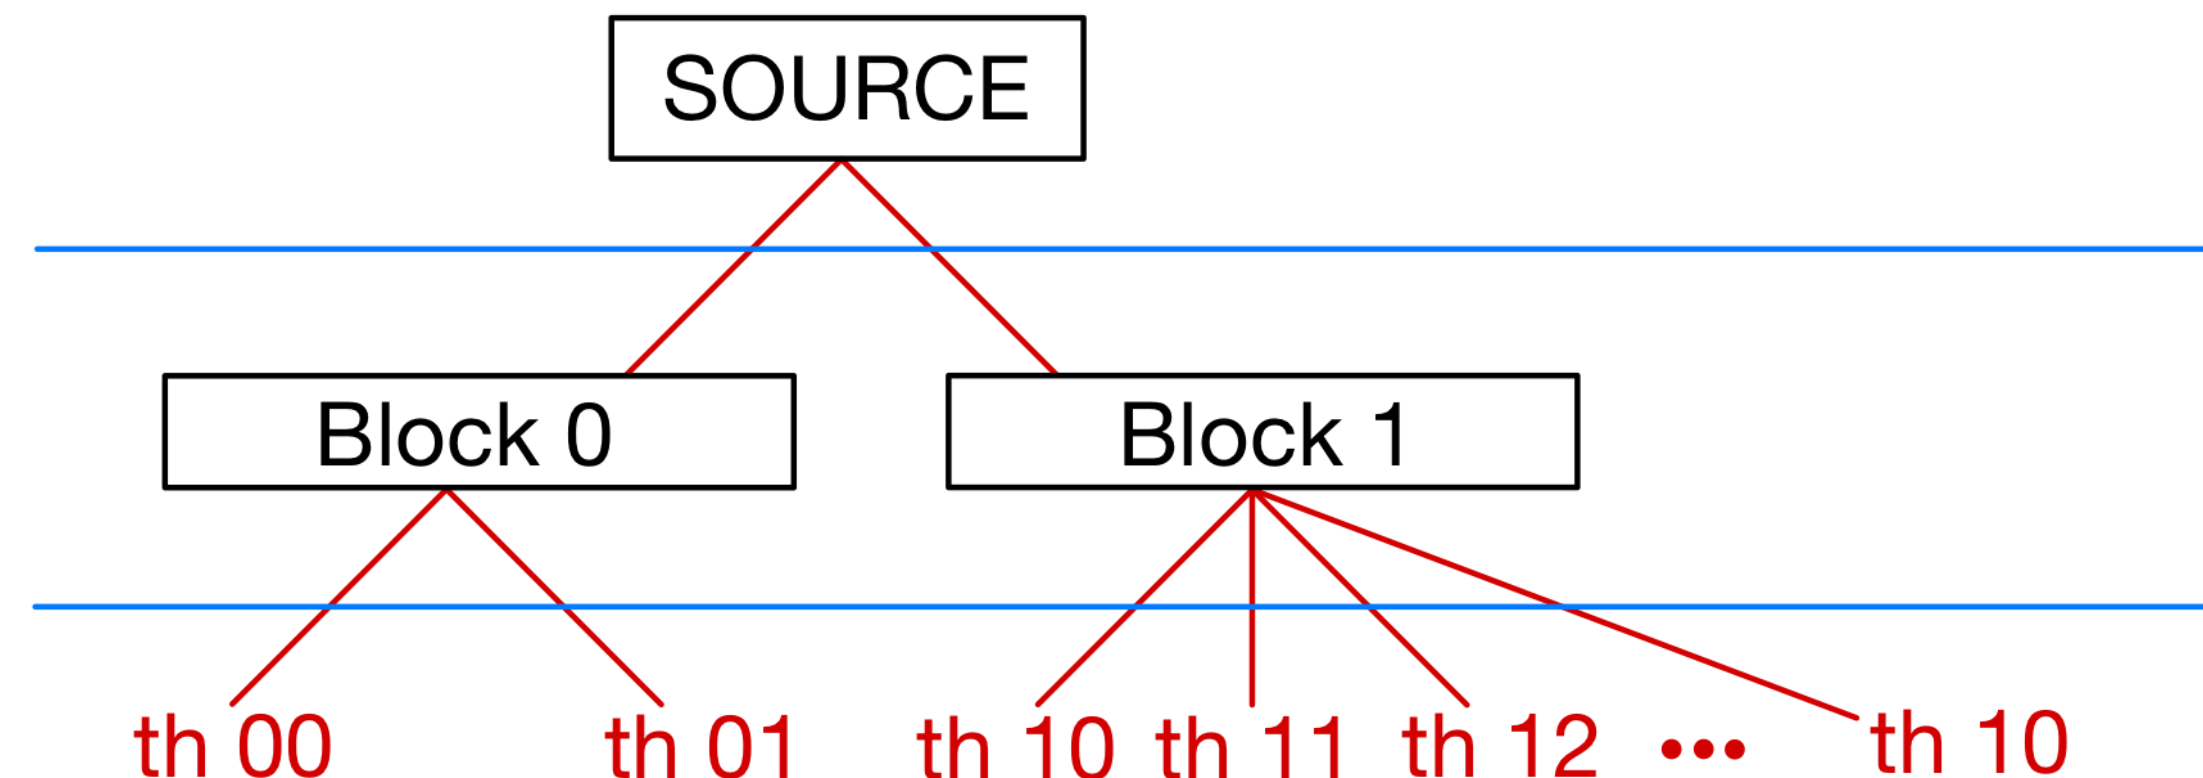
// Call the gpu_kernel function
BFS_gpu<<<DIM_GRID,DIM_BLOCK>>>(gpu_source, gpu_row_ptr, gpu_col_ind, gpu_dist);
CHECK_KERNELCALL();

// gpu memory free
CHECK(cudaFree(gpu_source));
CHECK(cudaFree(gpu_row_ptr));
CHECK(cudaFree(gpu_col_ind));
CHECK(cudaFree(gpu_dist));
```

# Procedure

## GPU architecture implementation

- Since the *BFS Algorithm* operates by Breadth Analysis, for each frontier are defined the following associations:
  - each vertex corresponds to a block of the *GPU*.
  - each edge of a vertex is analysed by a thread of the corresponding block.
  - if there are more edges in one vertex than thread per block, each thread may have to analyse sequentially more edges.





# Procedure

## kernell call

- The GPU-kernell call contains four variables:
  - a pointer to the root vertex const int \*source\_ptr.
  - a pointer to the first edge of each vertex const int \*rowPointers.
  - a pointer to the destination of each edge const int \*destinations.
  - a pointer wich contains the distances of all vertices from the root and is set at **-1** from the **main()** int \*distances.

```
// BFS algorithm optimized for GPU
__global__ void BFS_gpu(const int *source_ptr, const int *rowPointers, const int *destinations, int *distances)
{
```

# Procedure

## frontiers initialization

- For each block are initialised four `__shared__` variables:
  - an array representing the previous frontier of the Tree `int previousFrontier[ ]`.
  - an array representing the current frontier of the Tree `int currentFrontier[ ]`.
  - two integers used to indicate the size of the two frontiers (the number of outgoing edges from the frontier) `int previousFrontierSize` and `int currentFrontierSize`.

```
// initialize frontiers
__shared__ int currentFrontier[MAX_FRONTIER_SIZE];
__shared__ int currentFrontierSize;
__shared__ int previousFrontier[MAX_FRONTIER_SIZE];
__shared__ int previousFrontierSize;
```

# Procedure

## BFS loop setting

- Given the pointer to the root `const int *source_ptr`, the following step is to define the first frontier as `previousFrontier[ ]` and the first size as `previousFrontierSize` via the `__device__` function ***insertIntoFrontier()***.
- Only then is possible to start a `while()` loop, which is the core of the *BFS*.
- This loop will go on until the given frontier `previousFrontier[ ]` has at least a new child vertex to analyse ( `previousFrontierSize > 0` ).

# Procedure

## BFS loop setting

```
// initialize block's previous frontier from source
if (threadIdx.x == 0)
{
    currentFrontierSize = 0;
    previousFrontierSize = 0;
    const int source = *source_ptr;
    insertIntoFrontier(source, previousFrontier, &previousFrontierSize);
    distances[source] = 0;
}

__syncthreads();

// BFS with parallel vertices
while(previousFrontierSize > 0)    // while there are new vertices to visit
{
```

# Procedure

## BFS loop explanation

- At each iteration is analysed a single frontier, wich is represented by the couple previousFrontier[ ] and previousFrontierSize.
- Each frontier use only the necessary number of blocks in order to represent all the vertices via the check condition: if ( blockIdx.x < previousFrontierSize).
- For each block are defined:
  - the current vertex as int currentVertex = previousFrontier[ blockIdx.x ].
  - the address of the first edge of the currentVertex, wich is defined as int row\_start = rowPointers[ currentVertex ].
  - the address of the first edge of the next vertex, wich is used in order to establish the limit for the currentVertex edges, as int row\_end = rowPointers[ currentVertex +1 ].



# Procedure

## BFS loop explanation

```
// BFS with parallel vertices
while(previousFrontierSize > 0)      // while there are new vertices to visit
{
    // visit all vertices on the previous frontier
    if(blockIdx.x < previousFrontierSize)
    {
        int currentVertex = previousFrontier[blockIdx.x];
        int row_start = rowPointers[currentVertex];
        int row_end = rowPointers[currentVertex + 1];
    }
}
```

# Procedure

## BFS loop explanation

- The threads are managed in a more organised way:
- For each vertex represented by a block:
  - each edge is analysed by a single thread, via the association between the edge `int row_i = row_start + threadIdx.x`
  - a thread represents one edge only if it is true that `row_i < row_end`.
  - if the edges of a vertex are more than the number of threads per block, some threads may have to analyse more edges.
  - in order to cover this eventuality, the check is part of a `for()` loop, which is defined as `for ( int row_i; row_i < row_end; row_i += DIM_BLOCK)`.

# Procedure

## BFS loop explanation

- Then is carried out a check for each `int row_i`, which analyse the distance from the `currentVertex` and the destination.
  - if the vertex has not been explored yet, the check gives as a result: `distances[ destinations[ row_i ] ] == -1`.
  - in this case the `currentFrontier[ ]` and the `currentFrontierSize` have to be updated with the new vertex `destinations[ row_i ]` via the `__device__` function ***insertIntoFrontier()***.
  - it also has to be calculated the distance of the new vertex as one more than the `currentVertex` distance, as `distances[ destinations[ row_i ] ] = distances[ currentVertex ] + 1`.

# Procedure

## BFS loop explanation

```
// check all outgoing edges
for(int row_i = row_start + threadIdx.x; row_i < row_end; row_i += DIM_BLOCK)    // parallelize over all outgoing edges even if they are more than the block size
{
    if(distances[destinations[row_i]] == -1)
    {
        // this vertex has not been visited yet
        insertIntoFrontier(destinations[row_i], currentFrontier, &currentFrontierSize);
        distances[destinations[row_i]] = distances[currentVertex] + 1;
    }
}
```

# Procedure

- Once all the vertices have been analysed, all the new distances are updated in the frontier currentFrontier[ ].
- When all threads of each block have ended the analysis, the first thread per block calls the \_\_device\_\_ function **swap( )** between previousFrontier[ ] and currentFrontier[ ].
- When all the **swap( )** executions have come to an end, the loop is ready to begin a new iteration on the next frontier.

## BFS loop explanation

```
// wait for all vertices to be visited
__syncthreads();

// swap to the next frontier
if(threadIdx.x == 0)
{
    swap(&currentFrontier, &previousFrontier);
    previousFrontierSize = currentFrontierSize;
    currentFrontierSize = 0;
}

// synchronize with the swap
__syncthreads();
}
```



# Conclusions

## BFS implementation and total parallelisation

- We have seen an implementation of the ***Breadth-First Search (BFS) Algorithm*** in *CUDA* programming language, in order to use the capability of a *GPU* device.
- This implementation is also totally optimised, since it parallelise each level of the *GPU* architecture:
  - each vertex of the frontier is analysed by a block in parallel.
  - each edge of the vertex is analysed by a thread in parallel.
- The only limitation of this implementation is definition of *BFS* as a *Breadth Analysis Algorithm*, wich implies the sequentiality of the frontier analysis.